

PERFORMANCE AND FINAL SUBMISSION PHRASE

Date	21 November 2023
Team ID	Team- 592184
Project Name	ASL - Alphabet image recognition
Maximum Marks	10 Marks

Model Performance Testing:

S.No.	Parameter	Values	Screenshot
1.	Metrics	<p>Classification Model: VGG16 model</p> <p>Confusion Matrix , Classification Report & Accuracy Scores-</p> <p>Training accuracy:- 94.98%</p> <p>Testing accuracy:- 95.03%</p>	<p>Confusion Matrix:-</p> <pre> tf.Tensor([[[578 1 0 0 2 0 0 0 0 0] [7 1 1 0 0 1 0 3 0 0] [0 585 0 0 1 1 0 0 0 1] [0 0 5 0 0 0 0 0 0 1] [0 1 585 0 0 0 0 0 0 0] [1 1 0 0 0 0 0 2 0 3] [0 6 0 584 0 2 0 1 0 0] [0 0 0 1 0 0 0 0 0 0] [12 6 0 0 565 1 0 0 5 0 0] [7 0 0 1 0 0 0 0 0 0] [0 2 0 3 1 590 0 0 0 0] [1 0 0 0 0 0 1 0 0 0] [3 0 0 0 0 0 572 5 2 6] [2 0 0 2 0 0 1 1 0 0] [0 0 0 0 0 0 5 589 0 1] [0 0 0 0 0 0 0 2 0 2] [1 6 0 1 4 0 0 568 3 0] [1 0 1 1 0 0 2 3 0 0] [0 0 0 0 0 0 0 3 6 584] [0 0 0 0 0 0 4 3 0 0] [0 12 0 2 0 0 0 0 2 544] [0 0 0 27 3 0 0 0 1 0] [2 0 0 0 0 0 2 0 5 3] [1 10 0 0 0 3 1 1 0 0] [0 0 0 1 0 0 0 0 0 0] [1 0 0 0 0 0 0 0 0 0] [4 0 2 0 0 0 0 1 1 0] [0 0 0 2 0 0 0 0 0 0] [7 0 0 0 0 0 1 1 0 1] [0 0 1 0 0 0 0 0 1 0] [1 0 1 0 0 0 0 0 0 1] [0 0 0 0 0 0 0 1 4 0] [0 0 0 0 0 0 0 1 4 0] [6 0 18 6 0 4 0 0 0 0] [3 0 0 0 0 0 0 0 0 0] [578 2 2 0 0 4 0 3 0 1] [0 0 0 0 0 0 0 0 1 0] [12 576 0 0 0 4 6 0 0 0] [0 0 0 0 0 0 0 0 0 0] [4 0 539 2 0 0 0 0 0 0] [0 0 0 0 0 0 0 2 0 12] [7 1 5 548 0 6 0 0 0 0] [0 0 0 0 0 0 0 0 6 0] [2 0 1 22 563 0 2 1 0 0] [2 0 0 0 0 0 0 1 0 0] [44 3 3 1 0 529 0 5 0 1] [4 0 0 0 0 0 590 2 0 0] [1 0 0 0 0 0 0 0 0 0] [9 0 0 0 0 1 3 584 0 0] [0 0 0 0 0 0 2 0 0 0] [1 0 0 0 0 0 1 2 584 2] [0 0 0 0 0 0 0 0 0 0] [0 1 0 0 0 0 0 0 599 0]]]) </pre> <p>Classification report:-</p>

			<pre> 136/136 [=====] - 25s 181ms/step precision recall f1-score support A 0.95 0.95 0.95 600 B 0.94 0.96 0.95 600 C 1.00 0.97 0.98 600 D 0.99 0.98 0.98 600 E 0.98 0.94 0.96 600 F 0.99 0.97 0.98 600 G 0.99 0.95 0.97 600 H 0.97 0.98 0.97 600 I 0.97 0.96 0.96 600 J 0.97 0.97 0.97 600 K 0.94 0.91 0.93 600 L 1.00 0.96 0.98 600 M 0.93 0.93 0.93 600 N 0.94 0.95 0.94 600 O 0.97 0.99 0.98 600 P 0.98 0.98 0.98 600 Q 0.99 0.98 0.98 600 R 0.84 0.93 0.88 600 S 0.81 0.96 0.88 600 T 0.98 0.95 0.96 600 U 0.91 0.89 0.90 600 V 0.91 0.91 0.91 600 W 0.98 0.93 0.96 600 X 0.96 0.87 0.91 600 Y 0.97 0.97 0.97 600 Z 0.95 0.98 0.96 600 del 0.98 0.97 0.98 600 nothing 0.98 1.00 0.99 600 space 0.98 0.98 0.98 600 accuracy 0.95 17400 macro avg 0.96 17400 weighted avg 0.96 17400 </pre>
			<p>Accuracy Score:-</p> <pre> Epoch 1/30 543/543 [=====] - ETA: 0s - loss: 2.6236 - accuracy: 0.1762 Epoch 1: val_accuracy improved from -inf to 0.41754, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 163s 273ms/step - loss: 2.6236 - accuracy: 0.1762 - val_loss: 1.6868 - val_accuracy: 0.4175 Epoch 2/30 543/543 [=====] - ETA: 0s - loss: 1.8716 - accuracy: 0.6257 Epoch 2: val_accuracy improved from 0.41754 to 0.78942, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 170s 239ms/step - loss: 1.8716 - accuracy: 0.6257 - val_loss: 0.6494 - val_accuracy: 0.7894 Epoch 3/30 543/543 [=====] - ETA: 0s - loss: 0.5376 - accuracy: 0.8276 Epoch 3: val_accuracy improved from 0.78942 to 0.84824, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 144s 205ms/step - loss: 0.5376 - accuracy: 0.8276 - val_loss: 0.4681 - val_accuracy: 0.8482 Epoch 4/30 543/543 [=====] - ETA: 0s - loss: 0.3992 - accuracy: 0.8792 Epoch 4: val_accuracy improved from 0.84824 to 0.89718, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 129s 237ms/step - loss: 0.3992 - accuracy: 0.8792 - val_loss: 0.3558 - val_accuracy: 0.8972 Epoch 5/30 543/543 [=====] - ETA: 0s - loss: 0.3364 - accuracy: 0.9027 Epoch 5: val_accuracy improved from 0.89718 to 0.90845, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 132s 243ms/step - loss: 0.3364 - accuracy: 0.9027 - val_loss: 0.3819 - val_accuracy: 0.9085 Epoch 6/30 543/543 [=====] - ETA: 0s - loss: 0.2839 - accuracy: 0.9199 Epoch 6: val_accuracy improved from 0.90845 to 0.92727, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 148s 272ms/step - loss: 0.2839 - accuracy: 0.9199 - val_loss: 0.2608 - val_accuracy: 0.9273 Epoch 7/30 543/543 [=====] - ETA: 0s - loss: 0.2397 - accuracy: 0.9331 Epoch 7: val_accuracy did not improve from 0.92727 543/543 [=====] - 131s 241ms/step - loss: 0.2397 - accuracy: 0.9331 - val_loss: 0.2851 - val_accuracy: 0.9226 Epoch 8/30 543/543 [=====] - ETA: 0s - loss: 0.2162 - accuracy: 0.9397 Epoch 8: val_accuracy improved from 0.92727 to 0.94934, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 129s 238ms/step - loss: 0.2162 - accuracy: 0.9397 - val_loss: 0.1797 - val_accuracy: 0.9493 Epoch 9/30 543/543 [=====] - ETA: 0s - loss: 0.2057 - accuracy: 0.9445 Epoch 9: val_accuracy improved from 0.94934 to 0.95282, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 142s 262ms/step - loss: 0.2057 - accuracy: 0.9445 - val_loss: 0.1645 - val_accuracy: 0.9526 Epoch 10/30 543/543 [=====] - ETA: 0s - loss: 0.1810 - accuracy: 0.9498 Epoch 10: val_accuracy improved from 0.95282 to 0.95789, saving model to /content/sample_data/best_model_weights.h5 543/543 [=====] - 141s 259ms/step - loss: 0.1810 - accuracy: 0.9498 - val_loss: 0.1669 - val_accuracy: 0.9578 scores = model.evaluate(test_generator) print("No: %2FNN" % ("Evaluate Test Accuracy", scores[1]*100)) 136/136 [=====] - 25s 182ms/step - loss: 0.1469 - accuracy: 0.9626 Evaluate Test Accuracy: 96.263688 </pre>

2.	Tune the Model	<p><u>Hyperparameter Tuning:</u></p> <p>The model is tuned with following hyper parameters</p> <p>Optimizer - Adam</p> <p>Learning rate - 0.0001</p> <p>Loss - Categorical cross entropy</p> <p>Batch size - 128</p> <p>EPOCHS - 10</p> <p><u>Validation Method:</u></p> <p>The validation of the model is done through the validation data , which is set to 20% of training data. Data augmentation and callbacks are also used to validate performance. Accuracy is the validation parameter that we have monitored</p>	<p>Hyperparameter Tuning:-</p> <pre> 1 # Compile the model 2 model.compile(optimizer=Adam(lr=0.0001), loss='categorical_crossentropy', metrics=['accuracy']) 3 </pre> <p>Configuration</p> <pre> class CFG: # Set the batch size for training batch_size = 128 # Set the height and width of input images img_height = 32 img_width = 32 epochs = 10 </pre> <p>Validation Method:-</p> <pre> # Split the data into train and test sets X_train, X_test, y_train, y_test = train_test_split(metadata['image_path'], metadata['label'], test_size=0.2, random_state=2253, shuffle=True, stratify=metadata['label']) # Create a DataFrame for the training set test set data_train = pd.DataFrame({ 'image_path': X_train, 'label': y_train }) data_test = pd.DataFrame({ 'image_path': X_test, 'label': y_test }) # Create a ModelCheckpoint callback checkpoint_callback = ModelCheckpoint(filepath='./content/sample_data/best_model_weights.h5', monitor='val_accuracy', # Monitor validation accuracy for saving the best model save_best_only=True, mode='max', verbose=1) </pre>
----	----------------	--	--

CONFUSION MATRIX:-

```
tf.Tensor(
[[578  1  0  0  2  0  0  0  0  0  0  0  5  1  0  0  0  0
  7  1  1  0  0  1  0  3  0  0  0  0]  0  0  0  0  0  0  0
 [ 0 585  0  0  1  1  0  0  0  0  1  0  0  0  1  0  0  5
  0  0  5  0  0  0  0  0  0  0  1  0]  0  0  0  0  0  0  0
 [ 0  1 585  0  0  0  0  3  0  0  0  0  0  0  3  0  1  0
  1  1  0  0  0  0  0  0  2  0  3]  0  0  0  0  0  0  0
 [ 0  6  0 584  0  2  0  0  1  0  0  0  0  0  5  0  0  1
  0  0  0  1  0  0  0  0  0  0]  0  0  0  0  0  0  0
 [12  6  0  0 565  1  0  0  5  0  0  0  0  2  0  1  0  0
  7  0  0  1  0  0  0  0  0  0  0]  0  0  0  0  0  0  0
 [ 0  2  0  3  1 590  0  0  0  0  0  0  0  1  0  1  0  0
  1  0  0  0  0  0  1  0  0  0  0]  0  0  0  0  0  0  0
 [ 3  0  0  0  0  0 572  5  2  6  0  0  0  0  0  5  0  0
  2  0  0  2  0  0  1  1  0  0  1]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  5 589  0  1  0  0  0  0  1  0  0  0
  0  0  0  0  0  0  0  2  0  0]  0  0  0  0  0  0  0
 [ 1  6  0  1  4  0  0  0 568  8  0  0  0  0  0  0  0  4
  1  0  1  1  0  0  2  3  0  0  0]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  3  6 584  0  0  0  0  0  0  0
  0  0  0  0  0  0  4  3  0  0]  0  0  0  0  0  0  0
 [ 0 12  0  2  0  0  0  0  2  0 544  0  0  0  0  0  0  9
  0  0  0 27  3  0  0  0  0  1  0]  0  0  0  0  0  0  0
 [ 2  0  0  0  0  0  2  0  5  3  0 572  0  0  0  0  0  0
  1 10  0  0  0  3  1  1  0  0  0]  0  0  0  0  0  0  0
 [ 0  0  0  0  1  0  0  0  0  0  0  0 578 18  2  0  0  0
  1  0  0  0  0  0  0  0  0  0  0]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  1  0  0  0  39 551  0  0  0  0
  4  0  3  0  0  0  0  0  1  1  0]  0  0  0  0  0  0  0
 [ 0  0  0  2  0  0  0  0  0  0  0  0  0  0 585  2  0  1
  7  0  0  0  0  0  0  0  1  1  0  1]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  0  0  1  0  0  0  0 593  2  0
  1  0  1  0  0  0  0  0  1  0  1]  0  0  0  0  0  0  0
 [ 0  0  1  0  0  0  0  0  0  0  0  0  0  0  0  1 592  0
  0  0  0  0  0  0  0  1  4  0  1]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  1  4  0  1  0  0  0  0  0 560
  6  0 18  6  0  4  0  0  0  0]  0  0  0  0  0  0  0
 [ 3  0  0  0  0  0  0  0  0  0  0  0  0  3  3  0  0  0
 578  2  2  0  0  4  0  3  0  1  1]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  0  0  1  0  0  1  0  0  0  0
 12 576  0  0  0  0  4  6  0  0]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  1  0  0 50
  4  0 539  2  0  4  0  0  0]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  0  0  2  0 12  0  1  0  0  10
  7  1  5 548  8  6  0  0  0]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  0  0  0  0  6  0  1  0  0  2
  2  0  1 22 563  0  2  1  0]  0  0  0  0  0  0  0
 [ 2  0  0  0  0  0  0  0  1  0  0  0  1  0  0  0  0  8
 44  3  3  1  0 529  0  5  0  1]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  0  0  3  0  0  0  0  0  0  0
  4  1  0  0  0  0 590  2  0  0]  0  0  0  0  0  0  0
 [ 1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  1  0  1
  9  0  0  0  0  1  3 584  0  0]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  2  0  0  0  0  1  1  1  2  0
  1  0  0  0  0  0  1  2 584  2]  0  0  0  0  0  0  0
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  1  0  0  0  0  0  0 599  0]  0  0  0  0  0  0  0
```

CLASSIFICATION REPORT:-

136/136 [=====] - 25s 181ms/step

	precision	recall	f1-score	support
A	0.95	0.95	0.95	600
B	0.94	0.96	0.95	600
C	1.00	0.97	0.98	600
D	0.99	0.98	0.98	600
E	0.98	0.94	0.96	600
F	0.99	0.97	0.98	600
G	0.99	0.95	0.97	600
H	0.97	0.98	0.97	600
I	0.97	0.96	0.96	600
J	0.97	0.97	0.97	600
K	0.94	0.91	0.93	600
L	1.00	0.96	0.98	600
M	0.93	0.93	0.93	600
N	0.94	0.95	0.94	600
O	0.97	0.99	0.98	600
P	0.98	0.98	0.98	600
Q	0.99	0.98	0.98	600
R	0.84	0.93	0.88	600
S	0.81	0.96	0.88	600
T	0.98	0.95	0.96	600
U	0.91	0.89	0.90	600
V	0.91	0.91	0.91	600
W	0.98	0.93	0.96	600
X	0.96	0.87	0.91	600
Y	0.97	0.97	0.97	600
Z	0.95	0.98	0.96	600
del	0.98	0.97	0.98	600
nothing	0.98	1.00	0.99	600
space	0.98	0.98	0.98	600
accuracy			0.95	17400
macro avg	0.96	0.95	0.95	17400
weighted avg	0.96	0.95	0.95	17400

ACCURACY SCORE:-

⚠ WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.Adam.

```
Epoch 1/10
543/543 [=====] - ETA: 0s - loss: 2.5409 - accuracy: 0.1969
Epoch 1: val_accuracy improved from -inf to 0.42445, saving model to /content/sample_data/best_model_weights.h5
543/543 [=====] - 163s 269ms/step - loss: 2.5409 - accuracy: 0.1969 - val_loss: 1.7111 - val_accuracy: 0.4244
Epoch 2/10
543/543 [=====] - ETA: 0s - loss: 1.1009 - accuracy: 0.6191
Epoch 2: val_accuracy improved from 0.42445 to 0.73911, saving model to /content/sample_data/best_model_weights.h5
543/543 [=====] - 142s 261ms/step - loss: 1.1009 - accuracy: 0.6191 - val_loss: 0.7840 - val_accuracy: 0.7391
Epoch 3/10
543/543 [=====] - ETA: 0s - loss: 0.5650 - accuracy: 0.8141
Epoch 3: val_accuracy improved from 0.73911 to 0.87142, saving model to /content/sample_data/best_model_weights.h5
543/543 [=====] - 125s 230ms/step - loss: 0.5650 - accuracy: 0.8141 - val_loss: 0.4190 - val_accuracy: 0.8714
Epoch 4/10
543/543 [=====] - ETA: 0s - loss: 0.3615 - accuracy: 0.8875
Epoch 4: val_accuracy did not improve from 0.87142
543/543 [=====] - 143s 263ms/step - loss: 0.3615 - accuracy: 0.8875 - val_loss: 0.4847 - val_accuracy: 0.8604
Epoch 5/10
543/543 [=====] - ETA: 0s - loss: 0.3016 - accuracy: 0.9099
Epoch 5: val_accuracy improved from 0.87142 to 0.93473, saving model to /content/sample_data/best_model_weights.h5
543/543 [=====] - 125s 230ms/step - loss: 0.3016 - accuracy: 0.9099 - val_loss: 0.2324 - val_accuracy: 0.9347
Epoch 6/10
543/543 [=====] - ETA: 0s - loss: 0.2637 - accuracy: 0.9230
Epoch 6: val_accuracy did not improve from 0.93473
543/543 [=====] - 143s 263ms/step - loss: 0.2637 - accuracy: 0.9230 - val_loss: 0.2255 - val_accuracy: 0.9334
Epoch 7/10
543/543 [=====] - ETA: 0s - loss: 0.2280 - accuracy: 0.9352
Epoch 7: val_accuracy improved from 0.93473 to 0.94279, saving model to /content/sample_data/best_model_weights.h5
543/543 [=====] - 147s 270ms/step - loss: 0.2280 - accuracy: 0.9352 - val_loss: 0.2000 - val_accuracy: 0.9428
Epoch 8/10
543/543 [=====] - ETA: 0s - loss: 0.2205 - accuracy: 0.9377
Epoch 8: val_accuracy did not improve from 0.94279
543/543 [=====] - 129s 237ms/step - loss: 0.2205 - accuracy: 0.9377 - val_loss: 0.2174 - val_accuracy: 0.9416
Epoch 9/10
543/543 [=====] - ETA: 0s - loss: 0.2054 - accuracy: 0.9439
Epoch 9: val_accuracy improved from 0.94279 to 0.95756, saving model to /content/sample_data/best_model_weights.h5
543/543 [=====] - 127s 234ms/step - loss: 0.2054 - accuracy: 0.9439 - val_loss: 0.1446 - val_accuracy: 0.9576
Epoch 10/10
543/543 [=====] - ETA: 0s - loss: 0.1876 - accuracy: 0.9472
Epoch 10: val_accuracy did not improve from 0.95756
543/543 [=====] - 128s 235ms/step - loss: 0.1876 - accuracy: 0.9472 - val_loss: 0.1809 - val_accuracy: 0.9490
```

HYPERPARAMETER TUNING:-

```
1 # Compile the model
2 model.compile(optimizer=Adam(lr=0.0001), loss='categorical_crossentropy', metrics=['accuracy'])
3
```

VALIDATION METHOD:-

```

▶ # Configuration
class CFG:
    # Set the batch size for training
    batch_size = 128
    # Set the height and width of input images
    img_height = 32
    img_width = 32
    epochs = 10

22 # Split the training set into training and validation sets
23 X_train, X_val, y_train, y_val = train_test_split(
24     data_train['image_path'],
25     data_train['label'],
26     test_size=0.2/0.7, # Assuming you want 20% for validation out of the training set
27     random_state=2253,
28     shuffle=True,
29     stratify=data_train['label']
30 )
31
32 # Create a DataFrame for the validation set
33 data_val = pd.DataFrame({
34     'image_path': X_val,
35     'label': y_val
36 })

4 # Create a ModelCheckpoint callback
5 checkpoint_callback = ModelCheckpoint(
6     filepath='/content/sample_data/best_model_weights.h5',
7     monitor='val_accuracy', # Monitor validation accuracy for saving the best model
8     save_best_only=True,
9     mode='max',
10    verbose=1
11 )

```